

Intro to Side-Channel Analysis: Concepts and Techniques

(AI-assisted) side-channel attacks on real-world crypto implementations

Hardwear.io 2024 (Amsterdam) training, Lecture 1, Day 1

Lejla Batina¹, Łukasz Chmielewski² & Péter Horváth¹

¹ Digital Security group, Radboud University, The Netherlands

² Centre for Research on Cryptography and Security, Masaryk University, Czechia

October 21, 2024

Training schedule

- Day 1: Intro to Side-channel attacks
 - Lecture: Side-channel attacks on crypto implementations
 - Two Assignments
- Day 2: Advanced attacks
 - Lecture: Leakage evaluation: TVLA and alternatives + Assignment
 - Lecture: Profiling attacks + Assignment
 - Lecture: AI basics
- Day 3: AI and SCA
 - Lecture: AI-assisted SCA
 - Assignments
 - Lecture: Leakage simulators
 - Lecture: Higher order attacks (if there is time)
 - Unfinished Exercises

Day 1

- Let's know each other!
- Basic concepts and principles – gray box, platforms, setups, side channels
- Side-channel leakage modeling
- Simple power analysis (SPA)
- Differential power analysis (DPA)
- Correlation power analysis (CPA)
- Assignments 1-2: Hands-on DPA attacks on software and hardware implementations
 - For now ignore the TVLA part
- Resources: <https://tinyurl.com/ysnznaka>¹

¹Full: <https://www.dropbox.com/scl/fo/bd9r3lvzbk7eilqgytroo/ADK3b5aUicY2hzshgYCCQ8g?rlkey=5uzcpfdvbnlh4z6zal63tqc&st=j5a1094h>

Outline

- 1 Introduction
- 2 Timing side-channel
- 3 Power consumption as a side channel
- 4 SPA
- 5 Differential Power Analysis (DPA)
- 6 EM side channel
- 7 Other side channels
- 8 CPA

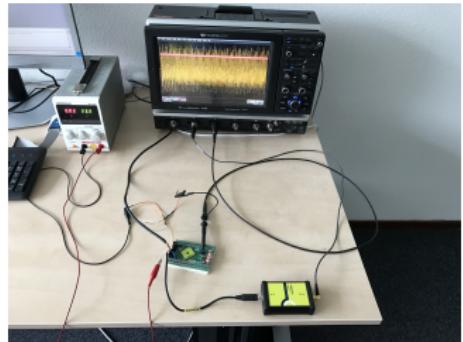
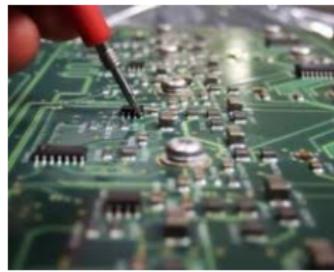
Introducing the side channel



Known challenge: embedded crypto devices



Implementation attacks



Relevance

Minerva, October 3, 2019

Researchers Discover ECDSA
Key Recovery Method



Relevance

TPM-FAIL, November 13, 2019



Minerva, October 3, 2019

Researchers Discover ECDSA
Key Recovery Method



Relevance

TPM-FAIL, November 13, 2019



LadderLeak, May 28, 2020

LadderLeak: Side-channel security flaws exploited to break ECDSA cryptography



Minerva, October 3, 2019

Researchers Discover ECDSA Key Recovery Method



Relevance

TPM-FAIL, November 13, 2019



LadderLeak, May 28, 2020

LadderLeak: Side-channel security flaws exploited to break ECDSA cryptography



SCA Titan: January 7, 2021



Minerva, October 3, 2019

Researchers Discover ECDSA Key Recovery Method
October 3, 2019 - Minerva by Facebook



Relevance

TPM-FAIL, November 13, 2019



LadderLeak, May 28, 2020

LadderLeak: Side-channel security flaws exploited to break ECDSA cryptography



SCA Titan: January 7, 2021



Minerva, October 3, 2019

Researchers Discover ECDSA Key Recovery Method



TPMScan, March 12, 2024

TPMScan: A wide-scale study of security-relevant properties of TPM 2.0 chips

Project Overview | About | News | Contact | Help

Project Leader: [Hans-Peter Selz](#) (University of Regensburg, Chair Cryptology)

Advisory Board: [Andreas Drosos](#), [Johannes Blaß](#), [Thomas Ristenpart](#)

Office Staff:

- [Julia Kühn](#) (University of Regensburg, Chair Cryptology)
- [Natalie Schäfer](#) (University of Regensburg, Chair Cryptology)
- [Helmut Pfeifer](#) (University of Regensburg, Chair Cryptology)
- [Hilma Pfeifer](#) (University of Regensburg, Chair Cryptology)
- [Gerald Lautenbacher](#) (University of Regensburg, Chair Cryptology)
- [Jens Freudenthal](#) (University of Regensburg, Chair Cryptology)

Press Contact: [Jens Freudenthal](#) (University of Regensburg, Chair Cryptology)

Project Status: [Open](#) | Last Update: 2024-03-12

Project Description: TPMScan is a wide-scale study of security-relevant properties of TPM 2.0 chips. The project aims to identify and analyze various security vulnerabilities and weaknesses in the TPM 2.0 chip ecosystem. By conducting a comprehensive analysis, we hope to contribute to the improvement of security standards and practices in the field of TPM 2.0 chips.

Project URL: <https://tpmscan.cs.uni-regensburg.de/>

Report URL: <https://tpmscan.cs.uni-regensburg.de/reports/>



Usage Statistics

Using TPMScan, researchers from 10 countries have made significant contributions to the field.

Relevance

TPM-FAIL, November 13, 2019



LadderLeak, May 28, 2020

LadderLeak: Side-channel security flaws exploited to break ECDSA cryptography



SCA Titan: January 7, 2021



Minerva, October 3, 2019

Researchers Discover ECDSA Key Recovery Method



TPMScan, March 12, 2024

TPMScan: A wide-scale study of security-relevant properties of TPM 2.0 chips



EUCLEAK, Yesterday



Blackbox scenario



- Cryptographic function is a black box, parameterized with key, that maps plaintext into ciphertext

Blackbox scenario



- Cryptographic function is a black box, parameterized with key, that maps plaintext into ciphertext
- Analyzing the security in the blackbox scenario relates to classical cryptanalysis

Blackbox scenario



- Cryptographic function is a black box, parameterized with key, that maps plaintext into ciphertext
- Analyzing the security in the blackbox scenario relates to classical cryptanalysis
- Adversary's goal: secret key or plaintext recovery by observing plaintext/ciphertext pairs

Greybox scenario

- Crypto is **implemented on a real device** such as a microcontroller, FPGA, ASIC etc.

Greybox scenario

- Crypto is **implemented on a real device** such as a microcontroller, FPGA, ASIC etc.
- We can measure and process certain *physical quantities* in the device's vicinity

Greybox scenario

- Crypto is **implemented on a real device** such as a microcontroller, FPGA, ASIC etc.
- We can measure and process certain *physical quantities* in the device's vicinity
- Adversary's goal: secret key or plaintext recovery by observing plaintext/ciphertext pairs **and a side channel**

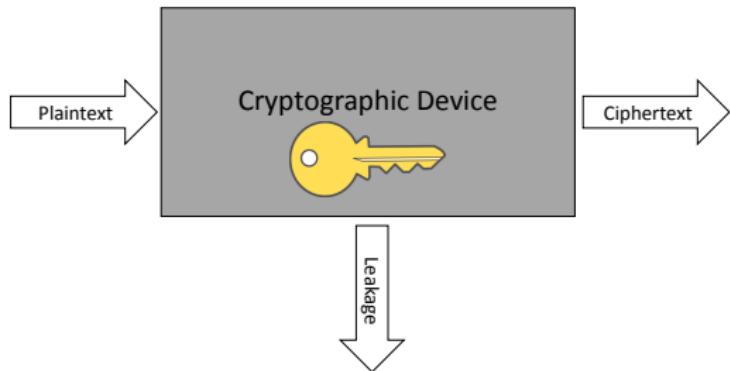
Greybox scenario

- Crypto is **implemented on a real device** such as a microcontroller, FPGA, ASIC etc.
- We can measure and process certain *physical quantities* in the device's vicinity
- Adversary's goal: secret key or plaintext recovery by observing plaintext/ciphertext pairs **and a side channel**
- *Side channel* is any unintentional signal that can offer us a blurry view of the algorithm's internal computations

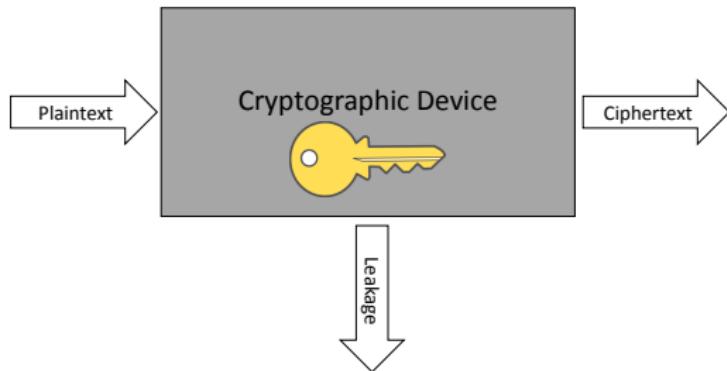
Greybox scenario

- Crypto is **implemented on a real device** such as a microcontroller, FPGA, ASIC etc.
- We can measure and process certain *physical quantities* in the device's vicinity
- Adversary's goal: secret key or plaintext recovery by observing plaintext/ciphertext pairs **and a side channel**
- *Side channel* is any unintentional signal that can offer us a blurry view of the algorithm's internal computations
- Examples: execution/reaction time, power consumption, electromagnetic emission, sound

Greybox scenario

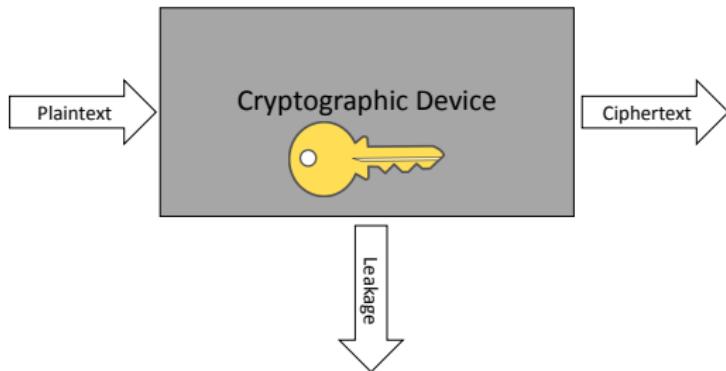


Greybox scenario



- Assuming limited access to the internal computations through this side channel window → greybox scenario

Greybox scenario



- Assuming limited access to the internal computations through this side channel window → greybox scenario
- Security in the blackbox scenario does not imply security under the greybox scenario

Timing side-channel



Timing side-channel: PIN verification

- Software for PIN code verification

Input: 4-digit PIN code

Output: PIN verified or rejected

Process CheckPIN (pin[4])

```
int pin_ok=0;
if (pin[0]==5)
    if (pin[1]==9)
        if (pin[2]==0)
            if (pin[3]==2)
                pin_ok=1;
            end
        end
    end
end
return pin_ok;
EndProcess
```

Timing side-channel: PIN verification

- Software for PIN code verification

```
Input: 4-digit PIN code
Output: PIN verified or rejected
Process CheckPIN (pin[4])
int pin_ok=0;
if (pin[0]==5)
    if (pin[1]==9)
        if (pin[2]==0)
            if (pin[3]==2)
                pin_ok=1;
            end
        end
    end
end
return pin_ok;
EndProcess
```

- What are the execution times of the process for PIN inputs

[0,1,2,3], [5,3,0,2], [5,9,0,0]

Timing side-channel: PIN verification

- Software for PIN code verification

```
Input: 4-digit PIN code
Output: PIN verified or rejected
Process CheckPIN (pin[4])
int pin_ok=0;
if (pin[0]==5)
    if (pin[1]==9)
        if (pin[2]==0)
            if (pin[3]==2)
                pin_ok=1;
            end
        end
    end
end
return pin_ok;
EndProcess
```

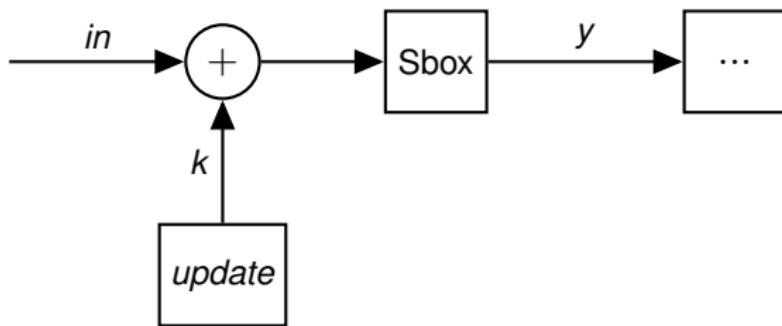
- What are the execution times of the process for PIN inputs

[0,1,2,3], [5,3,0,2], [5,9,0,0]

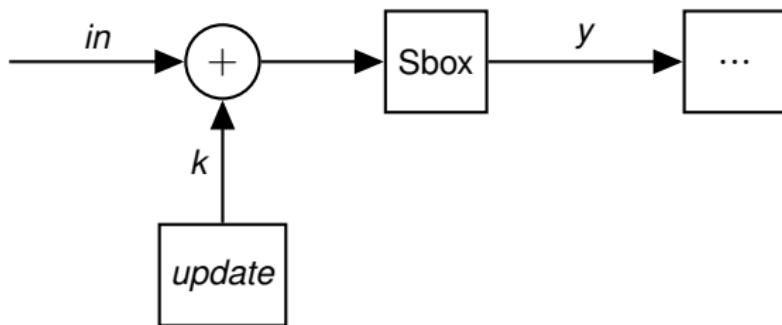
- The execution time increases as we get closer to

[5,9,0,2]

Timing side-channel: Cache attacks

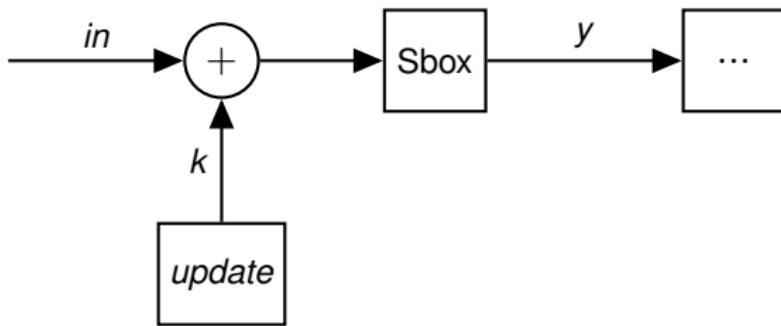


Timing side-channel: Cache attacks



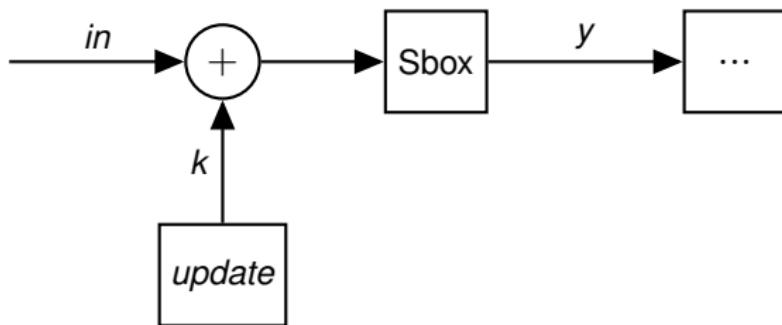
- Assume AES Sbox with a lookup table (LUT) stored in memory

Timing side-channel: Cache attacks



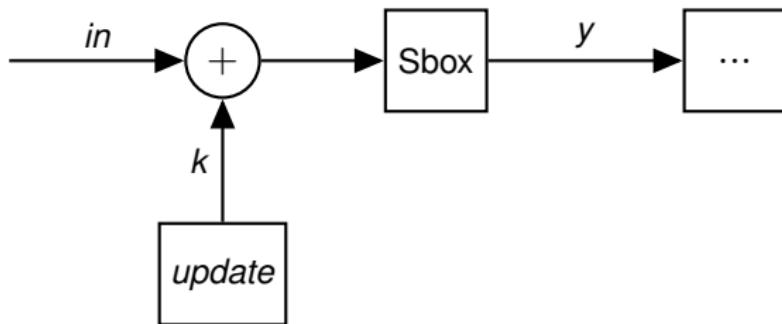
- Assume AES Sbox with a lookup table (LUT) stored in memory
- Every Sbox iteration creates a table index $j = \text{input} \oplus \text{key}$ and uses it to lookup the Sbox output, i.e. $y = \text{LUT}(j)$

Timing side-channel: Cache attacks



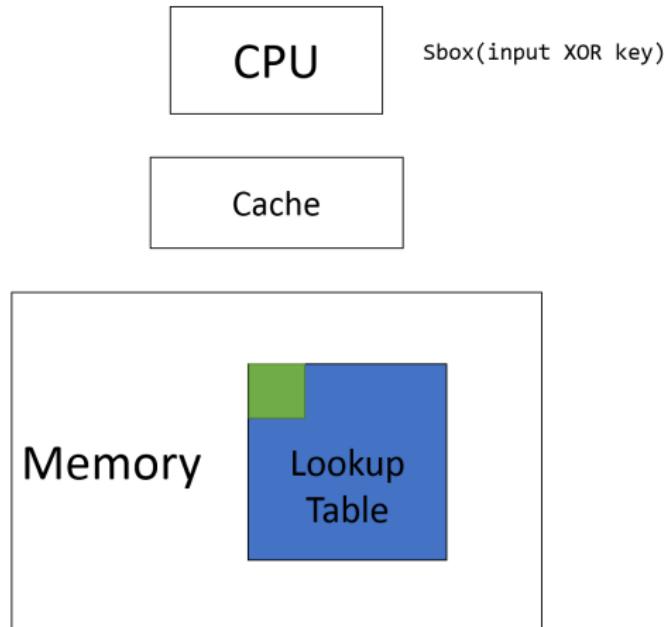
- Assume AES Sbox with a lookup table (LUT) stored in memory
- Every Sbox iteration creates a table index $j = \text{input} \oplus \text{key}$ and uses it to lookup the Sbox output, i.e. $y = \text{LUT}(j)$
- Accessing different parts of the lookup table may take different amount of time!

Timing side-channel: Cache attacks

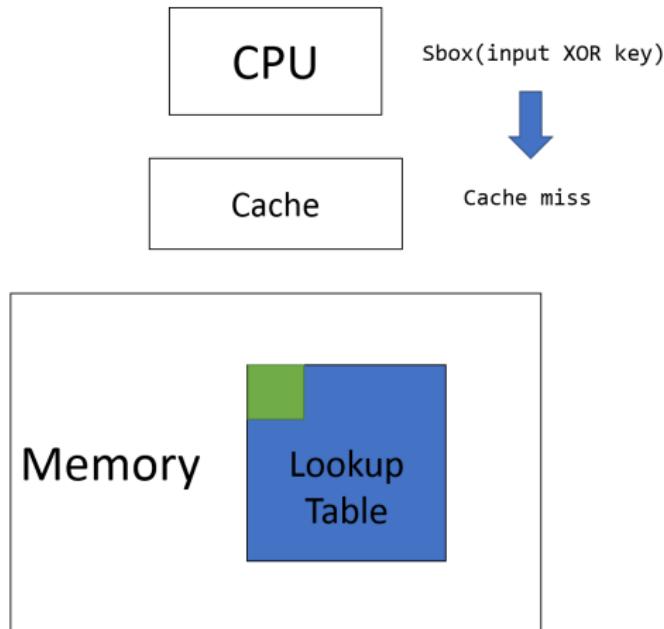


- Assume AES Sbox with a lookup table (LUT) stored in memory
- Every Sbox iteration creates a table index $j = \text{input} \oplus \text{key}$ and uses it to lookup the Sbox output, i.e. $y = \text{LUT}(j)$
- Accessing different parts of the lookup table may take different amount of time!
- The underlying cause is cache behavior of modern processors

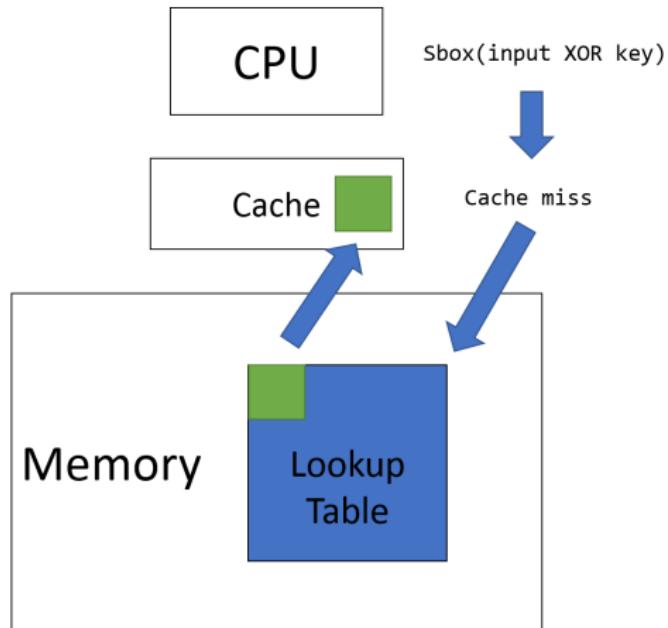
Timing side-channel: Cache attacks



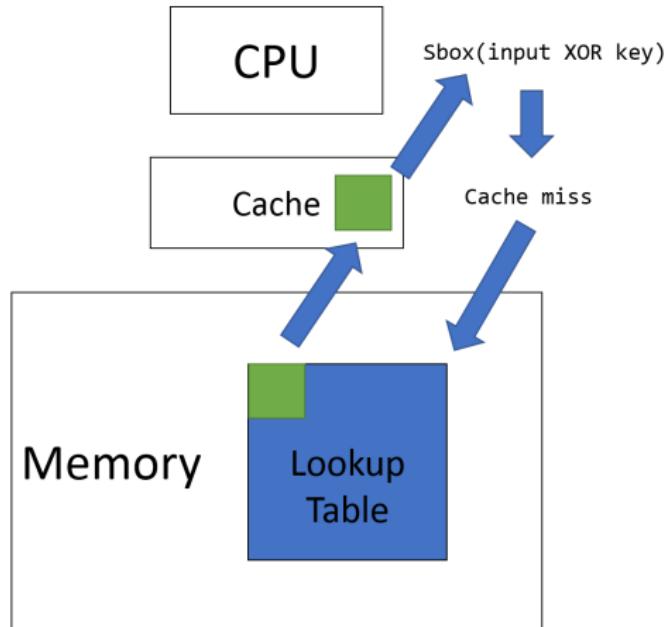
Timing side-channel: Cache behavior



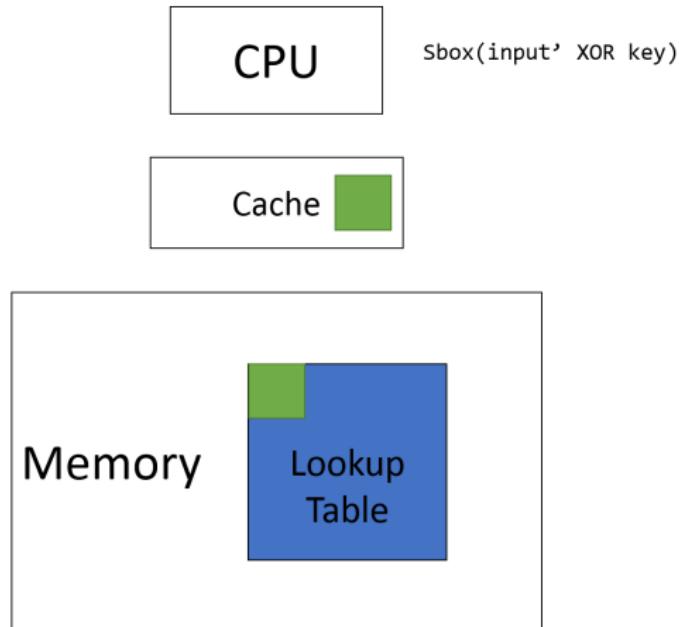
Timing side-channel: Cache behavior



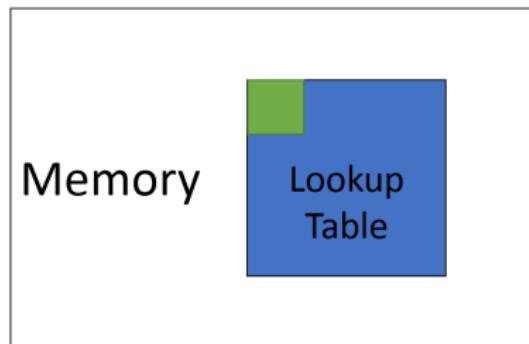
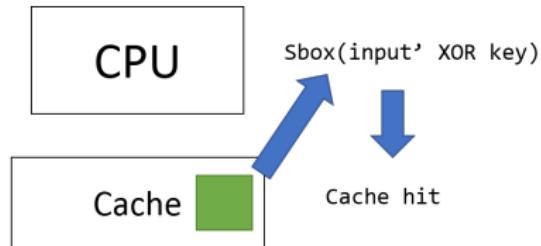
Timing side-channel: Cache behavior



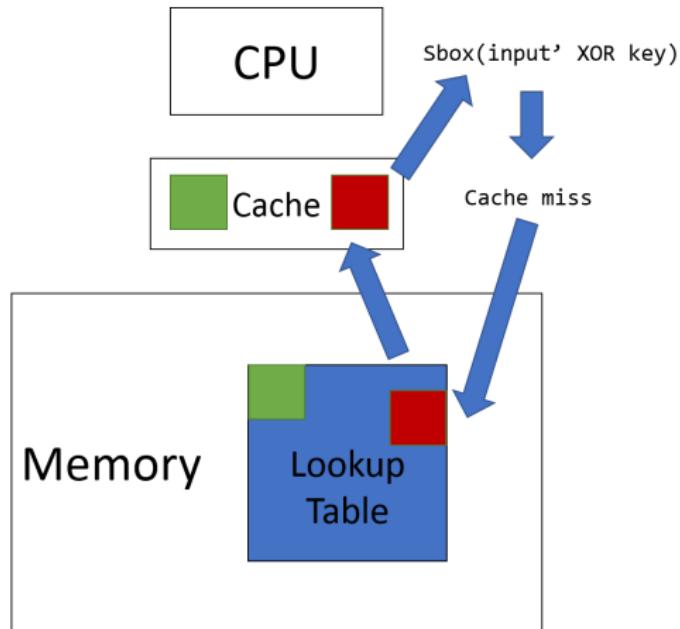
Timing side-channel: Cache behavior



Timing side-channel: Cache behavior



Timing side-channel: Cache behavior



Timing attacks notes and literature

- One of the earliest side-channel attacks due to easy measurements collection
- Can also be exploited remotely
- Exploits some not foreseen effects of caches to crypto implementations

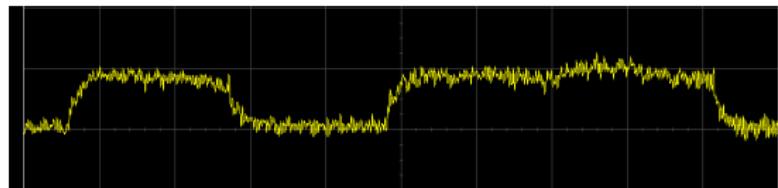
Timing attacks notes and literature

- One of the earliest side-channel attacks due to easy measurements collection
- Can also be exploited remotely
- Exploits some not foreseen effects of caches to crypto implementations
- Applied to symmetric and asymmetric cryptography

Timing attacks notes and literature

- One of the earliest side-channel attacks due to easy measurements collection
- Can also be exploited remotely
- Exploits some not foreseen effects of caches to crypto implementations
- Applied to symmetric and asymmetric cryptography
- P. Kocher. "Timing Attacks on Implementations of Diffie-Hellman, RSA, DSS, and Other Systems", CRYPTO 1996
- Daniel Page. "Theoretical use of cache memory as a cryptanalytic side-channel", technical report CSTR-02-003, 2002.
<https://eprint.iacr.org/2002/169.pdf>
- Daniel J. Bernstein. "Cache-timing attacks on AES."
<https://cr.yp.to/antiforgery/cachetiming-20050414.pdf>

Power side-channel



Power side channel: CMOS leakage

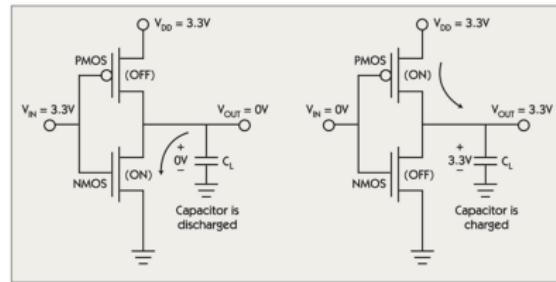


Fig. 1 Dynamic power in a CMOS inverter.

- CMOS is the most popular technology and the CMOS circuits exhibit several types of leakage

Power side channel: CMOS leakage

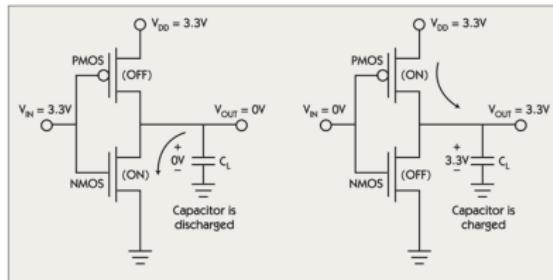


Fig. 1 Dynamic power in a CMOS Inverter.

- CMOS is the most popular technology and the CMOS circuits exhibit several types of leakage
- The most relevant for side-channel attacks is the charge and discharge of the CMOS load capacitance a.k.a dynamic power consumption

Power side channel: CMOS leakage

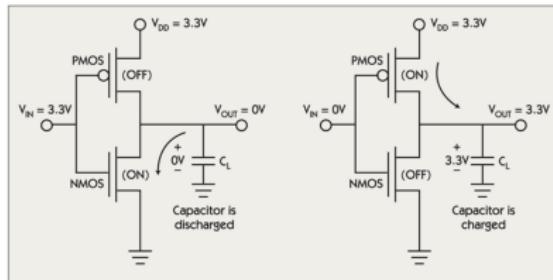


Fig. 1 Dynamic power in a CMOS Inverter.

- CMOS is the most popular technology and the CMOS circuits exhibit several types of leakage
- The most relevant for side-channel attacks is the charge and discharge of the CMOS load capacitance a.k.a dynamic power consumption
- Power analysis attack exploits the fact that the dynamic power consumption depends on the data and instructions being processed

Power side channel: CMOS leakage

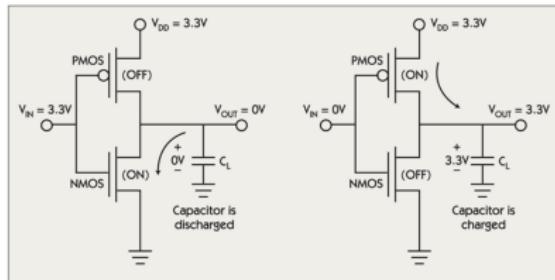


Fig. 1 Dynamic power in a CMOS inverter.

- CMOS is the most popular technology and the CMOS circuits exhibit several types of leakage
- The most relevant for side-channel attacks is the charge and discharge of the CMOS load capacitance a.k.a dynamic power consumption
- Power analysis attack exploits the fact that the dynamic power consumption depends on the data and instructions being processed
- Dynamic power consumption (P_{dyn}) is produced by CMOS transitions from state 0 to 1 and from state 1 to 0

Power side channel: CMOS leakage

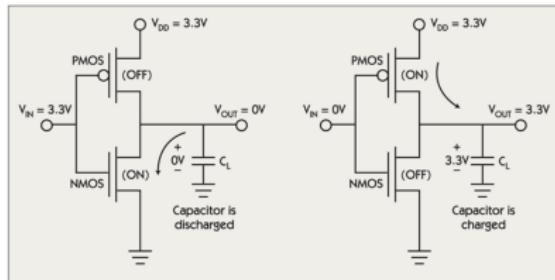


Fig. 1 Dynamic power in a CMOS Inverter.

- CMOS is the most popular technology and the CMOS circuits exhibit several types of leakage
- The most relevant for side-channel attacks is the charge and discharge of the CMOS load capacitance a.k.a dynamic power consumption
- Power analysis attack exploits the fact that the dynamic power consumption depends on the data and instructions being processed
- Dynamic power consumption (P_{dyn}) is produced by CMOS transitions from state 0 to 1 and from state 1 to 0
- $P_{dyn} = CV_{DD}^2 P_{0 \rightarrow 1} f$,
where C the transistor capacitance, V_{DD} the power supply voltage, f the frequency and $P_{0 \rightarrow 1}$ the probability of a $0 \rightarrow 1$ transition

Power side channel: CMOS leakage

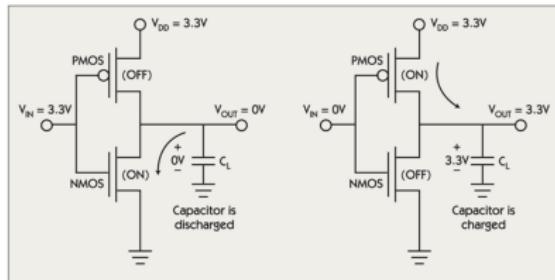


Fig. 1 Dynamic power in a CMOS Inverter.

- CMOS is the most popular technology and the CMOS circuits exhibit several types of leakage
- The most relevant for side-channel attacks is the charge and discharge of the CMOS load capacitance a.k.a dynamic power consumption
- Power analysis attack exploits the fact that the dynamic power consumption depends on the data and instructions being processed
- Dynamic power consumption (P_{dyn}) is produced by CMOS transitions from state 0 to 1 and from state 1 to 0
- $P_{dyn} = CV_{DD}^2 P_{0 \rightarrow 1} f$,
where C the transistor capacitance, V_{DD} the power supply voltage, f the frequency and $P_{0 \rightarrow 1}$ the probability of a $0 \rightarrow 1$ transition

Power side channel: Modeling the leakage

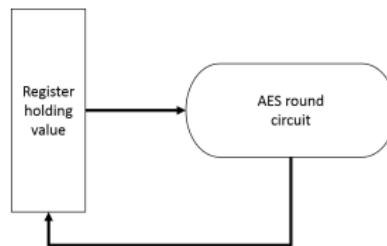
- Starting point: we use the number of transitions to model the leakage

Power side channel: Modeling the leakage

- Starting point: we use the number of transitions to model the leakage
- The Hamming distance model counts the number of $0 \rightarrow 1$ and $1 \rightarrow 0$ transitions

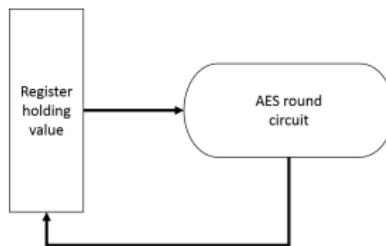
Power side channel: Modeling the leakage

- Starting point: we use the number of transitions to model the leakage
- The Hamming distance model counts the number of $0 \rightarrow 1$ and $1 \rightarrow 0$ transitions
- Example 1: A register R is storing the result of an AES round and initial value v_0 gets overwritten with v_1



Power side channel: Modeling the leakage

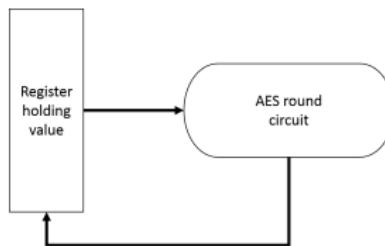
- Starting point: we use the number of transitions to model the leakage
- The Hamming distance model counts the number of $0 \rightarrow 1$ and $1 \rightarrow 0$ transitions
- Example 1: A register R is storing the result of an AES round and initial value v_0 gets overwritten with v_1



- The power consumption because of the register transition $v_0 \rightarrow v_1$ is related to the number of bit flips that occurred

Power side channel: Modeling the leakage

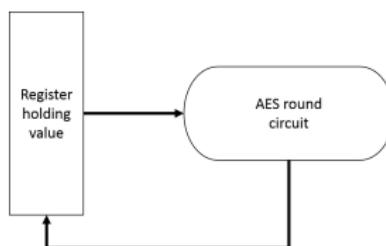
- Starting point: we use the number of transitions to model the leakage
- The Hamming distance model counts the number of $0 \rightarrow 1$ and $1 \rightarrow 0$ transitions
- Example 1: A register R is storing the result of an AES round and initial value v_0 gets overwritten with v_1



- The power consumption because of the register transition $v_0 \rightarrow v_1$ is related to the number of bit flips that occurred
- Thus it can be modeled as $\text{HammingDistance}(v_0, v_1) = \text{HammingWeight}(v_0 \oplus v_1)$

Power side channel: Modeling the leakage

- Starting point: we use the number of transitions to model the leakage
- The Hamming distance model counts the number of $0 \rightarrow 1$ and $1 \rightarrow 0$ transitions
- Example 1: A register R is storing the result of an AES round and initial value v_0 gets overwritten with v_1

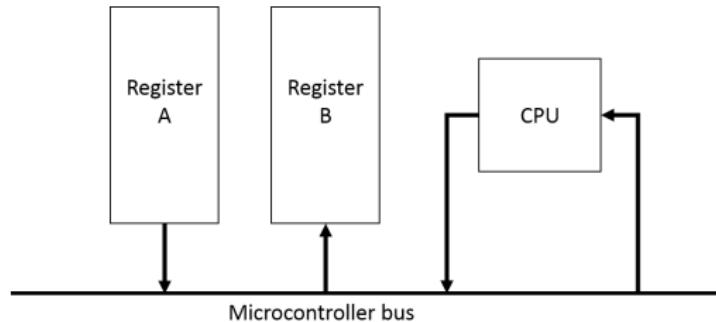


- The power consumption because of the register transition $v_0 \rightarrow v_1$ is related to the number of bit flips that occurred
- Thus it can be modeled as $\text{HammingDistance}(v_0, v_1) = \text{HammingWeight}(v_0 \oplus v_1)$
- Common leakage model for hardware implementations (FPGA, ASIC)

Power side channel: Modeling the leakage

- Example 2: In a microcontroller, a register A contains value v_0 and an assembly instruction moves the content of register A to B

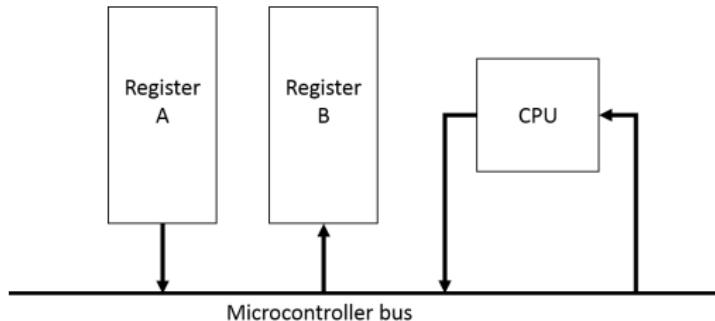
```
mov rB, rA
```



Power side channel: Modeling the leakage

- Example 2: In a microcontroller, a register A contains value v_0 and an assembly instruction moves the content of register A to B

```
mov rB, rA
```

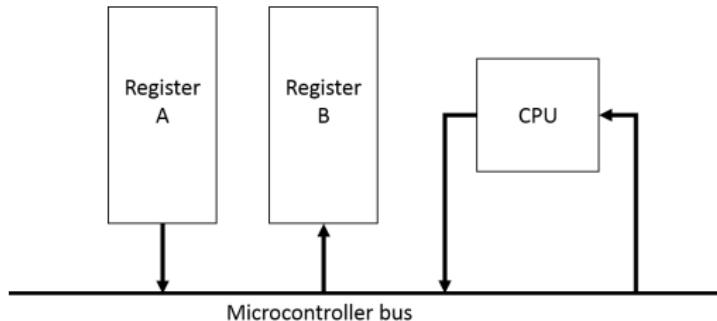


- This instruction transfers v_0 from A to B via the CPU, using the bus

Power side channel: Modeling the leakage

- Example 2: In a microcontroller, a register A contains value v_0 and an assembly instruction moves the content of register A to B

```
mov rB, rA
```

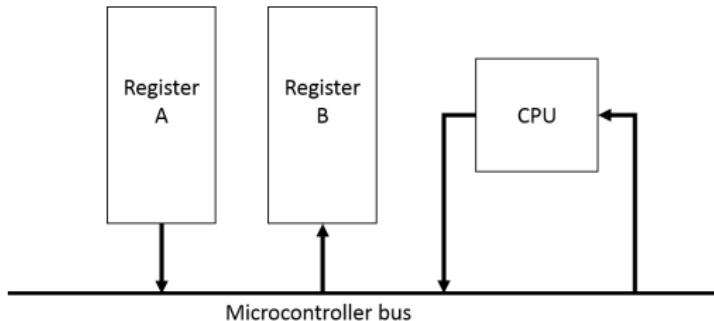


- This instruction transfers v_0 from A to B via the CPU, using the bus
- Typically the bus is precharged at all bits being zeros or one (busInitialValue)

Power side channel: Modeling the leakage

- Example 2: In a microcontroller, a register A contains value v_0 and an assembly instruction moves the content of register A to B

```
mov rB, rA
```

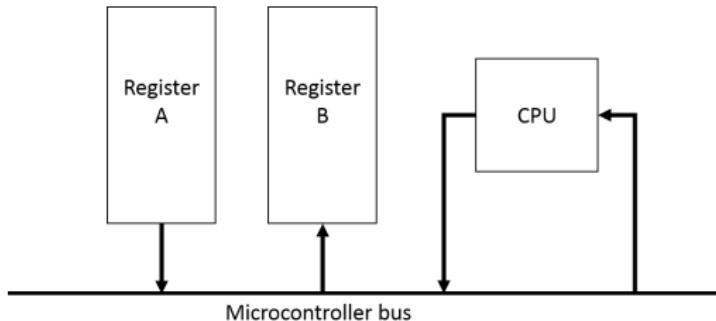


- This instruction transfers v_0 from A to B via the CPU, using the bus
- Typically the bus is precharged at all bits being zeros or one (`busInitialValue`)
- The power consumption of the instruction can be modeled as
 $\text{HammingDistance}(\text{busInitialValue}, v_0) = \text{HammingWeight}(v_0 \oplus 0) = \text{HW}(v_0)$

Power side channel: Modeling the leakage

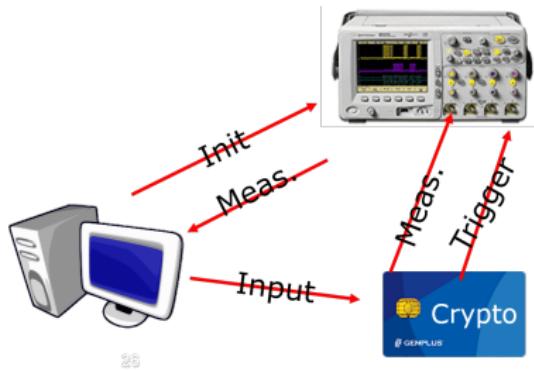
- Example 2: In a microcontroller, a register A contains value v_0 and an assembly instruction moves the content of register A to B

```
mov rB, rA
```



- This instruction transfers v_0 from A to B via the CPU, using the bus
- Typically the bus is precharged at all bits being zeros or one (`busInitialValue`)
- The power consumption of the instruction can be modeled as
 $\text{HammingDistance}(\text{busInitialValue}, v_0) = \text{HammingWeight}(v_0 \oplus 0) = \text{HW}(v_0)$
- Common leakage model for software implementations (AVR/ARM)

Measurement setup schematics

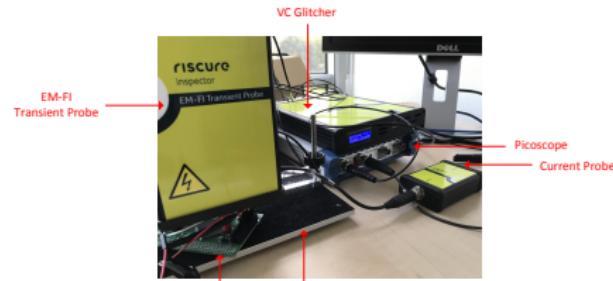


25

- Usually power measurements require physical proximity to the device and customized measurement equipment (resistor, oscilloscope)

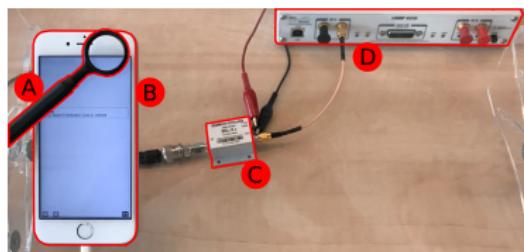
Actual setups

FA setup

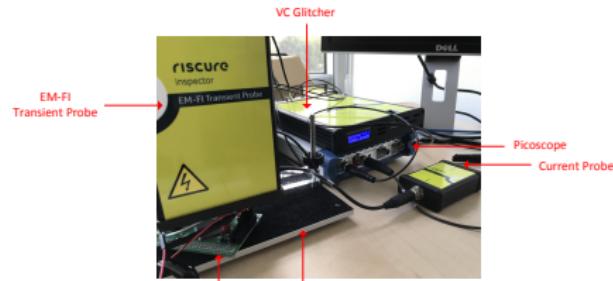


Actual setups

Tempest

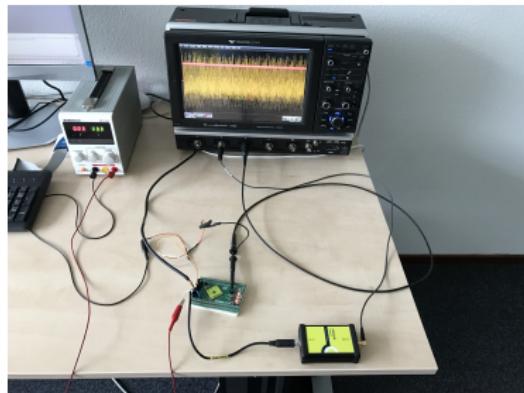


FA setup

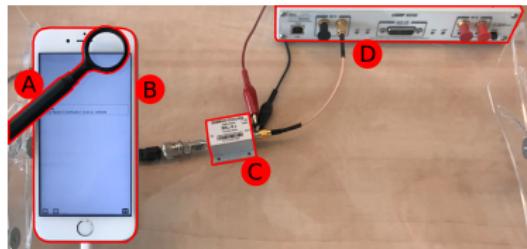


Actual setups

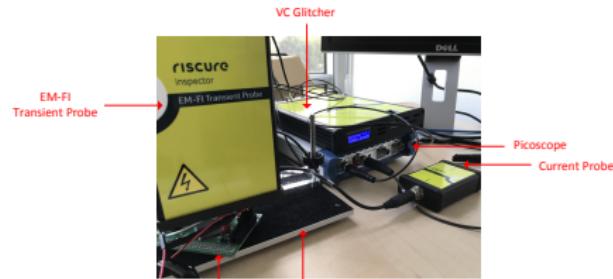
DPA setup with ARM CortexM4



Tempest

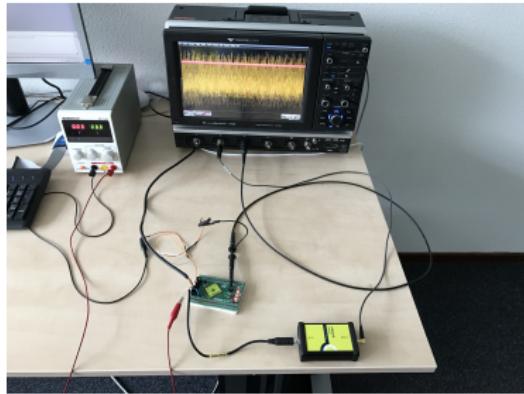


FA setup

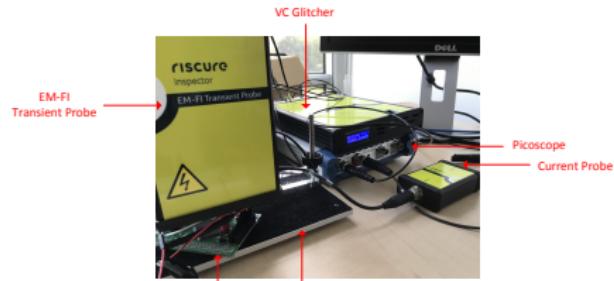


Actual setups

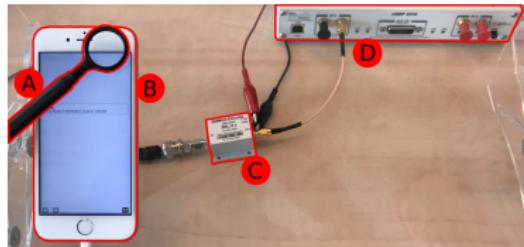
DPA setup with ARM CortexM4



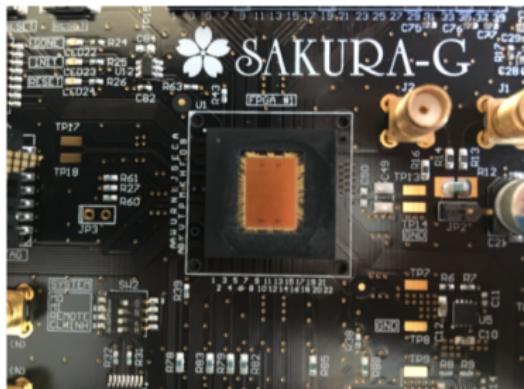
FA setup



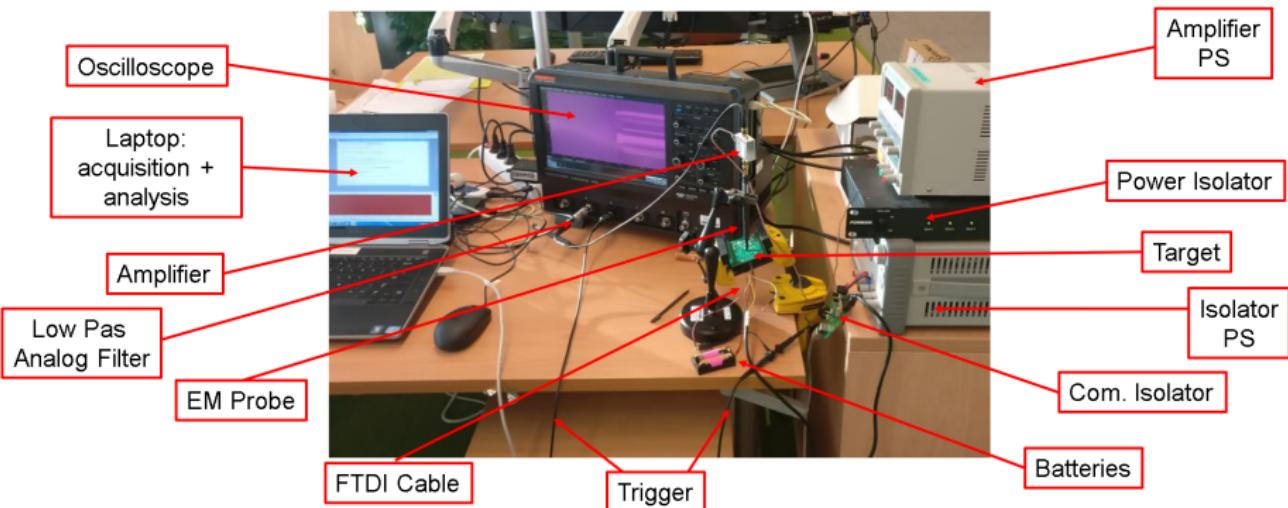
Tempest



FPGA board for SCA



Actual (overcomplicated?) setup



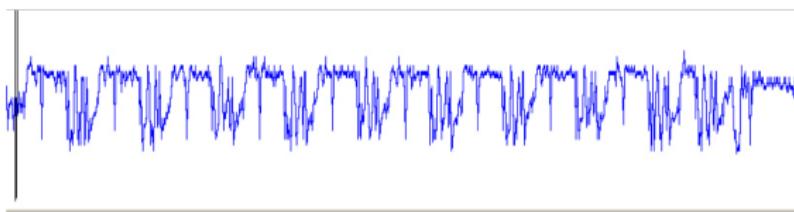
Analysis capabilities

- Simple Power Analysis (SPA): one or a few measurements - visual inspection or some simple signal processing
- Differential attacks (DPA): multiple measurements - use of statistics, signal processing, etc.
- Higher order attacks: Univariate vs multivariate
- Profiled attacks: more in the lecture on Template attacks
- Combining two or more side-channels
- Combining side-channel attack with theoretical cryptanalysis

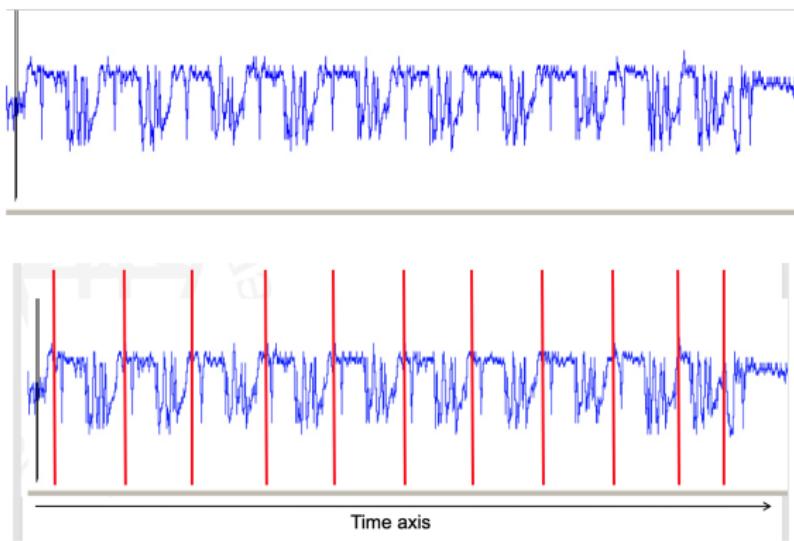
Simple Power Analysis (SPA)

- Based on one or a few measurements
- Mostly discovery of data-(in)dependent but instruction-dependent properties e.g.
 - Symmetric:
 - Number of rounds (resp. key length)
 - Memory accesses (usually higher power consumption)
 - Asymmetric:
 - The key (if badly implemented, e.g. RSA / ECC)
 - Key length
 - Implementation details: for example RSA w/wo CRT

SPA example 1



SPA example 1



SPA on RSA (Square-and-Multiply)

- RSA modular exponentiation

Input: integers x , e , n , length l of e

Output: $x^e \bmod n$

Process ModularExponentiation(x , e , n , l)

```
r=1;  
for j=l-1 down to 0  
    r=r^2 mod n //square  
    if (bit j of e) == 1  
        r= r*x mod n //multiply  
    end  
return r;  
EndProcess
```

SPA on RSA (Square-and-Multiply)

- RSA modular exponentiation

Input: integers x , e , n , length l of e

Output: $x^e \bmod n$

Process ModularExponentiation(x , e , n , l)

```
r=1;  
for j=l-1 down to 0  
    r=r^2 mod n //square  
    if (bit j of e) == 1  
        r= r*x mod n //multiply  
    end  
return r;  
EndProcess
```

- Do you already see a timing attack?

SPA on RSA (Square-and-Multiply)

- RSA modular exponentiation

Input: integers x , e , n , length l of e

Output: $x^e \bmod n$

Process ModularExponentiation(x , e , n , l)

```
r=1;  
for j=l-1 down to 0  
    r=r^2 mod n //square  
    if (bit j of e) == 1  
        r= r*x mod n //multiply  
    end  
return r;  
EndProcess
```

- Do you already see a timing attack?
- The exponent-dependent branch is causing it!

SPA on RSA (Square-and-Multiply)

- RSA modular exponentiation

Input: integers x , e , n , length l of e

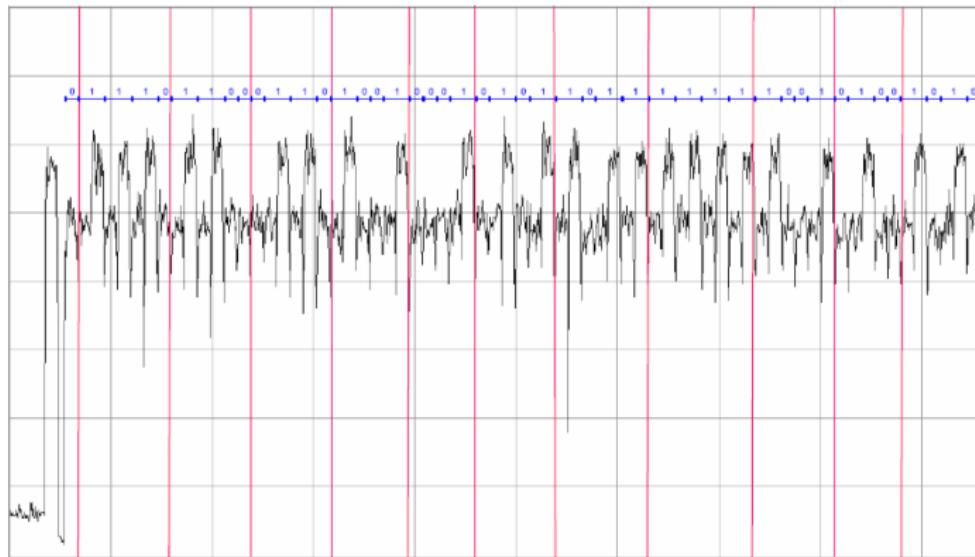
Output: $x^e \bmod n$

Process ModularExponentiation(x , e , n , l)

```
r=1;  
for j=l-1 down to 0  
    r=r^2 mod n //square  
    if (bit j of e) == 1  
        r= r*x mod n //multiply  
    end  
return r;  
EndProcess
```

- Do you already see a timing attack?
- The exponent-dependent branch is causing it!
- Do you see another side-channel attack?

SPA example 2: RSA exponentiation



DPA: main concepts

The DPA assumption: the attack is possible assuming the existence of a sensitive variable (for which exhaustive key search is possible) that depends on something we know (msg) and something we want to learn (key).

DPA: main concepts

The DPA assumption: the attack is possible assuming the existence of a sensitive variable (for which exhaustive key search is possible) that depends on something we know (msg) and something we want to learn (key).

Main steps

- 1 Choose your sensitive variable

DPA: main concepts

The DPA assumption: the attack is possible assuming the existence of a sensitive variable (for which exhaustive key search is possible) that depends on something we know (msg) and something we want to learn (key).

Main steps

- ① Choose your sensitive variable
- ② Collect measurements, known plaintext/ciphertext, sub-key guesses

DPA: main concepts

The DPA assumption: the attack is possible assuming the existence of a sensitive variable (for which exhaustive key search is possible) that depends on something we know (msg) and something we want to learn (key).

Main steps

- ① Choose your sensitive variable
- ② Collect measurements, known plaintext/ciphertext, sub-key guesses
- ③ Predict (hypothetical) intermediate values

DPA: main concepts

The DPA assumption: the attack is possible assuming the existence of a sensitive variable (for which exhaustive key search is possible) that depends on something we know (msg) and something we want to learn (key).

Main steps

- ① Choose your sensitive variable
- ② Collect measurements, known plaintext/ciphertext, sub-key guesses
- ③ Predict (hypothetical) intermediate values
- ④ Decide on the leakage model

DPA: main concepts

The DPA assumption: the attack is possible assuming the existence of a sensitive variable (for which exhaustive key search is possible) that depends on something we know (msg) and something we want to learn (key).

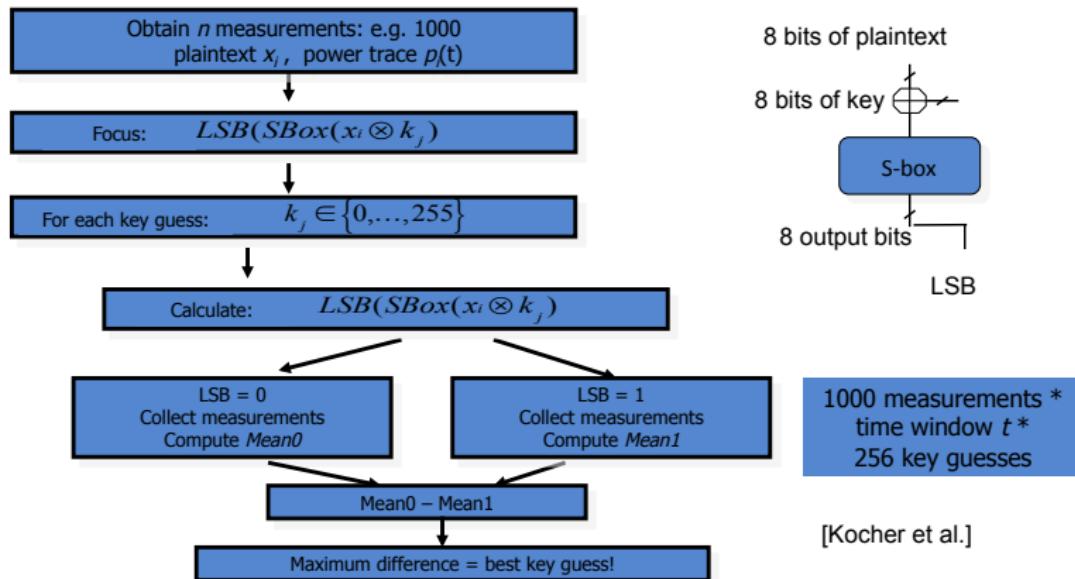
Main steps

- 1 Choose your sensitive variable
- 2 Collect measurements, known plaintext/ciphertext, sub-key guesses
- 3 Predict (hypothetical) intermediate values
- 4 Decide on the leakage model
- 5 Recover the key by statistical means, using partition or comparison method as the side-channel distinguisher

DPA with Distance of Means (DoM) as distinguisher

Classical 1-bit DPA on AES using DoM

AES impl.



Power analysis notes and literature

- Very powerful attacks that require contact with the target

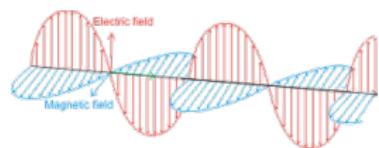
Power analysis notes and literature

- Very powerful attacks that require contact with the target
- Countermeasures on different layers required i.e. algorithm, implementation, transistor

Power analysis notes and literature

- Very powerful attacks that require contact with the target
- Countermeasures on different layers required i.e. algorithm, implementation, transistor
- P. Kocher, J. Jaffe, B. Jun. "Differential Power Analysis", CRYPTO 1999.
- T. Eisenbarth et al. "On the Power of Power Analysis in the Real World: A Complete Break of the KeeLoQCode Hopping Scheme", CRYPTO 2008.
- Mangard et al. Power Analysis Attacks, Springer, 2006.
- T. Kasper et al. "All You Can Eat or Breaking a Real-World Contactless Payment System", Financial Cryptography 2010.
- J. Balasch et al. "Power Analysis of Atmel CryptoMemory - Recovering Keys from Secure EEPROMs", CT-RSA 2012.
- N. Samwel et al. "Breaking Ed25519 in WolfSSL.", CT-RSA 2018.

Electromagnetic side channel



EM side channel: Probing

- Observing a power signal in more complex systems can be messy

EM side channel: Probing

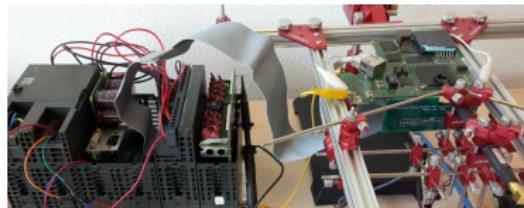
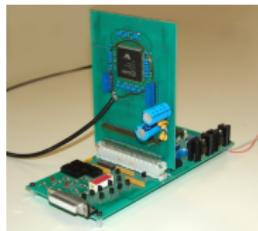
- Observing a power signal in more complex systems can be messy
- Complicated SoCs with multiple peripherals

EM side channel: Probing

- Observing a power signal in more complex systems can be messy
- Complicated SoCs with multiple peripherals
- Countermeasures trying to flatten the power consumption signal

EM side channel: Probing

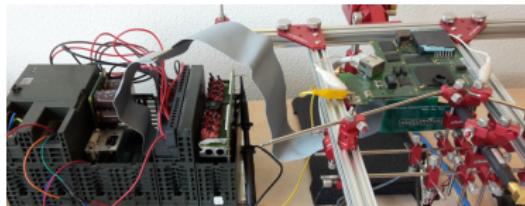
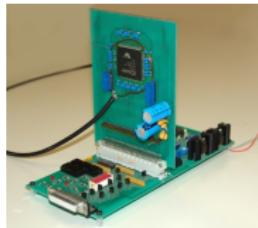
- Observing a power signal in more complex systems can be messy
- Complicated SoCs with multiple peripherals
- Countermeasures trying to flatten the power consumption signal
- Use an electromagnetic probe instead



- A probe is used to access the power consumption with less board modifications

EM side channel: Probing

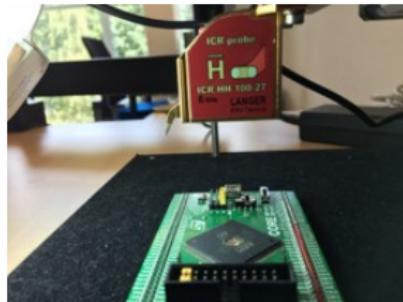
- Observing a power signal in more complex systems can be messy
- Complicated SoCs with multiple peripherals
- Countermeasures trying to flatten the power consumption signal
- Use an electromagnetic probe instead



- A probe is used to access the power consumption with less board modifications
- Smaller probes can focus on interesting locations and ignore interference from unrelated el. components

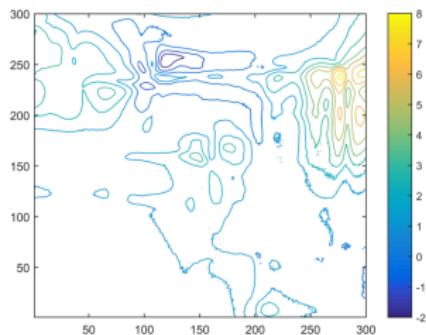
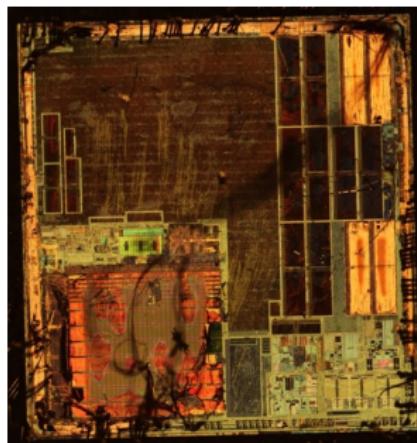
EM side channel: Decapsulation and Microprobing

- To improve spatial resolution of analysis use a micrometer-sized antenna
- To exploit more leakage decapsulate the chip using chemicals



EM side channel: Decapsulation and Microprobing

- Left: close inspection of decapsulated ARM processor using a microscope
- Right: EM emission heat-map of the same chip

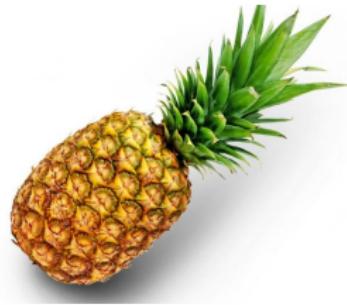


EM side channel: notes and literature

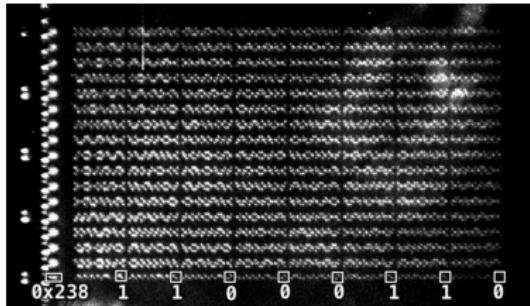
- EM enables side-channel attacks both in high proximity scenarios and distance scenarios
- Main side channel for SoCs, FPGAs, contactless cards due to their complexity and communication methods
- TEMPEST-like attacks are also targeting private data and authentication methods such as code etc.
- Gandolfi et al.: Electromagnetic Analysis: Concrete Results, CHES 1999.
- Andrikos et al.: Location, location, location: Revisiting modeling and exploitation for location-based side channel leakages, ASIACRYPT2019

Exotic side channels

Exotic side channels

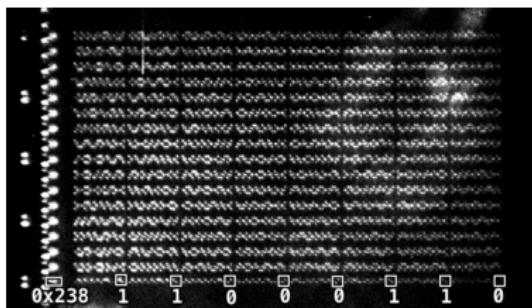


Optical Emission



- Accessing the chip SRAM cells emits photons that can be detected by a high-resolution camera
- Visual inspection can reveal the memory location accessed

Optical Emission



- Accessing the chip SRAM cells emits photons that can be detected by a high-resolution camera
- Visual inspection can reveal the memory location accessed
- The memory location maps to a specific value (e.g. in the AES LUT), i.e. it maps directly to $Sbox(in \oplus key)$
- Since the input in is known, knowledge of the memory location reveals the key
- Schlösser et al.: Simple Photonic Emission Analysis of AES, CHES 2012.

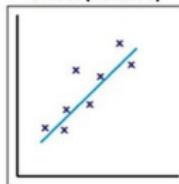
Out-of-order and speculative execution



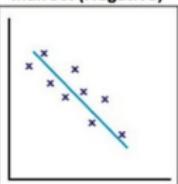
- Meltdown and Spectre Attacks
- Kernel addresses were accessed unintentionally due to out-of-order execution
- Taking all possible branches may also cause issues
- Foreshadow as a variant of Meltdown on SGX
- Recent ones: RIDL, ZombieLoad, ...

Correlation Power Analysis (CPA)

Direct (Positive)



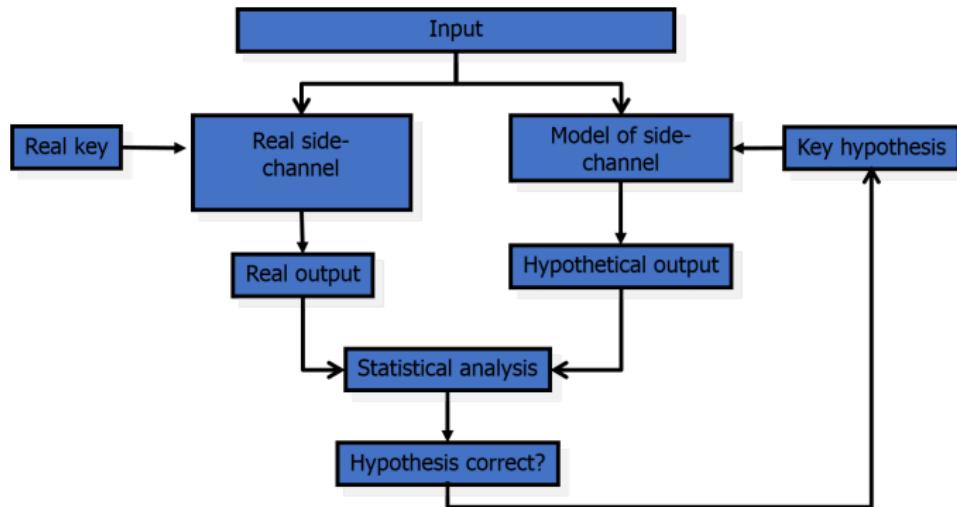
Indirect (Negative)



No Correlation

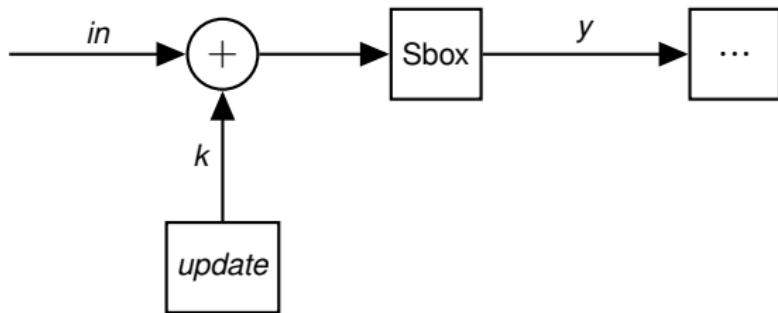


CPA: the main principle



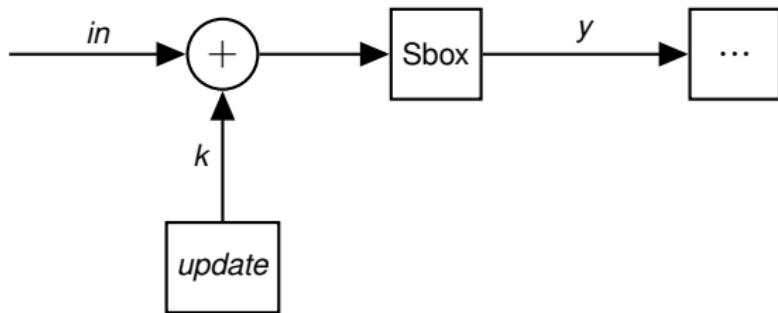
- Brier et al.: Correlation Power Analysis with a Leakage Model, CHES 2004

CPA Step 1: Choose intermediate value and decide on the leakage model



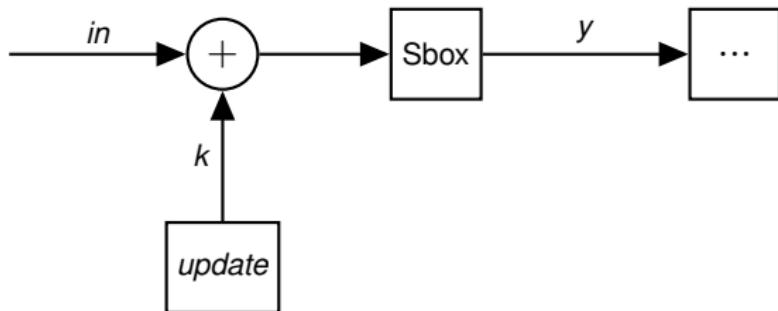
- Assume a software AES implementation e.g. for AVR microcontrollers (8-bit arch.)

CPA Step 1: Choose intermediate value and decide on the leakage model



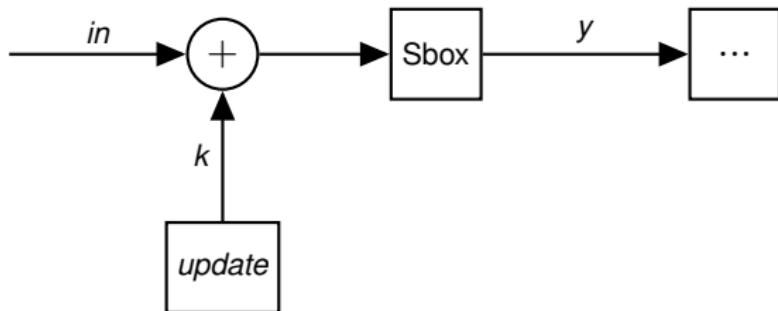
- Assume a software AES implementation e.g. for AVR microcontrollers (8-bit arch.)
- Step 1: Choose an intermediate value v of the AES cipher to attack

CPA Step 1: Choose intermediate value and decide on the leakage model



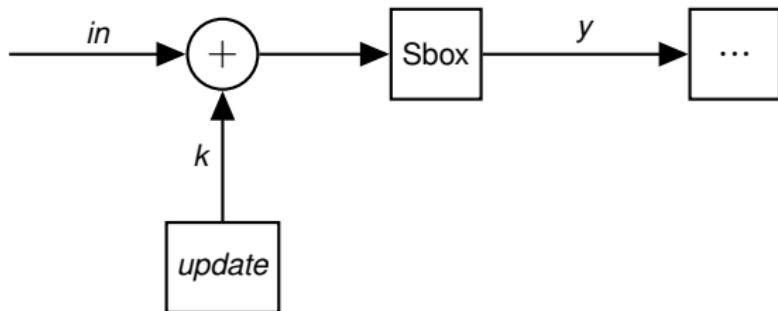
- Assume a software AES implementation e.g. for AVR microcontrollers (8-bit arch.)
- Step 1: Choose an intermediate value v of the AES cipher to attack
- The value v must be a function of the input and the key, i.e. $v = f(in, k)$

CPA Step 1: Choose intermediate value and decide on the leakage model



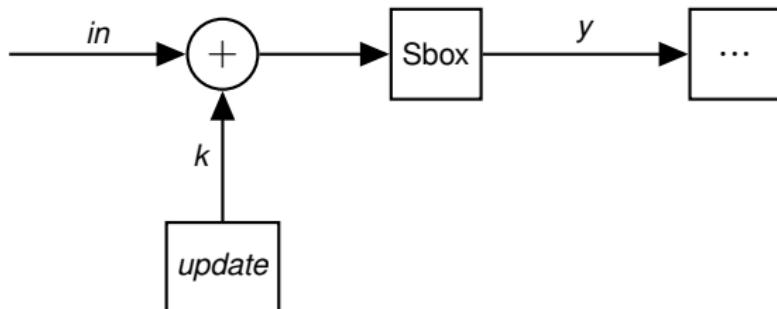
- Assume a software AES implementation e.g. for AVR microcontrollers (8-bit arch.)
- Step 1: Choose an intermediate value v of the AES cipher to attack
- The value v must be a function of the input and the key, i.e. $v = f(in, k)$
- A common choice for v is the Sbox output, i.e. $v = y = Sbox(in \oplus k)$

CPA Step 1: Choose intermediate value and decide on the leakage model



- Assume a software AES implementation e.g. for AVR microcontrollers (8-bit arch.)
- Step 1: Choose an intermediate value v of the AES cipher to attack
- The value v must be a function of the input and the key, i.e. $v = f(in, k)$
- A common choice for v is the Sbox output, i.e. $v = y = Sbox(in \oplus k)$
- Throughout the attack the key k must remain constant

CPA Step 1: Choose intermediate value and decide on the leakage model



- Assume a software AES implementation e.g. for AVR microcontrollers (8-bit arch.)
- Step 1: Choose an intermediate value v of the AES cipher to attack
- The value v must be a function of the input and the key, i.e. $v = f(in, k)$
- A common choice for v is the Sbox output, i.e. $v = y = Sbox(in \oplus k)$
- Throughout the attack the key k must remain constant
- Throughout the attack the input in is random

CPA Step 2: Measure the power consumption

- In Step 2 we record power consumption traces for multiple random inputs

CPA Step 2: Measure the power consumption

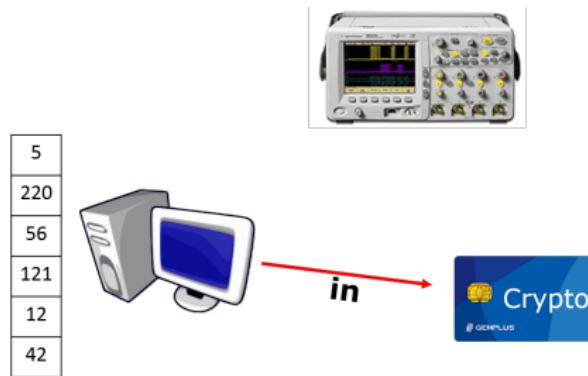
- In Step 2 we record power consumption traces for multiple random inputs
- Generate randomly n 8-bit inputs. Typically n is large (thousands to millions!)

CPA Step 2: Measure the power consumption

- In Step 2 we record power consumption traces for multiple random inputs
- Generate randomly n 8-bit inputs. Typically n is large (thousands to millions!)
- Store the inputs in vector $\mathbf{in} = [in_1 \ in_2 \ in_3 \dots \ in_n]^T$

CPA Step 2: Measure the power consumption

- In Step 2 we record power consumption traces for multiple random inputs
- Generate randomly n 8-bit inputs. Typically n is large (thousands to millions!)
- Store the inputs in vector $\mathbf{in} = [in_1 \ in_2 \ in_3 \dots \ in_n]^T$



CPA Step 2: Measure the power consumption

- For every generated 8-bit input we measure the power consumption of the AES implementation over time

CPA Step 2: Measure the power consumption

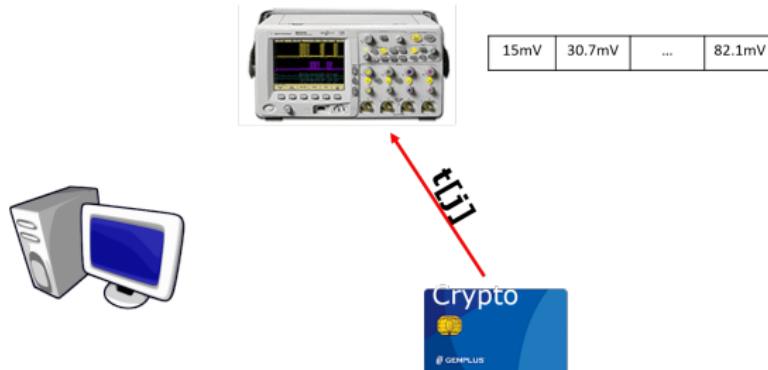
- For every generated 8-bit input we measure the power consumption of the AES implementation over time
- For every input $in_j, j = 1, \dots, n$ we capture a digitized trace over time

CPA Step 2: Measure the power consumption

- For every generated 8-bit input we measure the power consumption of the AES implementation over time
- For every input $in_j, j = 1, \dots, n$ we capture a digitized trace over time
- We denote the trace related to input in_j as $\mathbf{t}_j = [t_j^1 t_j^2 \dots t_j^m]^T$. It contains m points in time (a.k.a. samples)

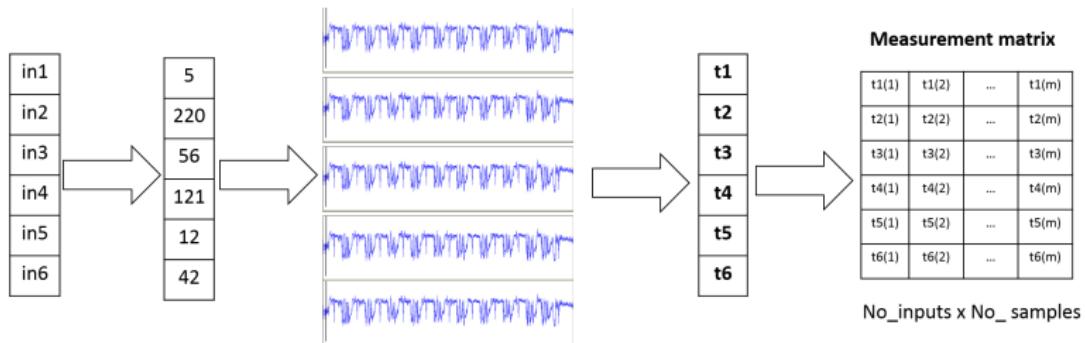
CPA Step 2: Measure the power consumption

- For every generated 8-bit input we measure the power consumption of the AES implementation over time
- For every input $in_j, j = 1, \dots, n$ we capture a digitized trace over time
- We denote the trace related to input in_j as $t_j = [t_j^1 t_j^2 \dots t_j^m]^T$. It contains m points in time (a.k.a. samples)



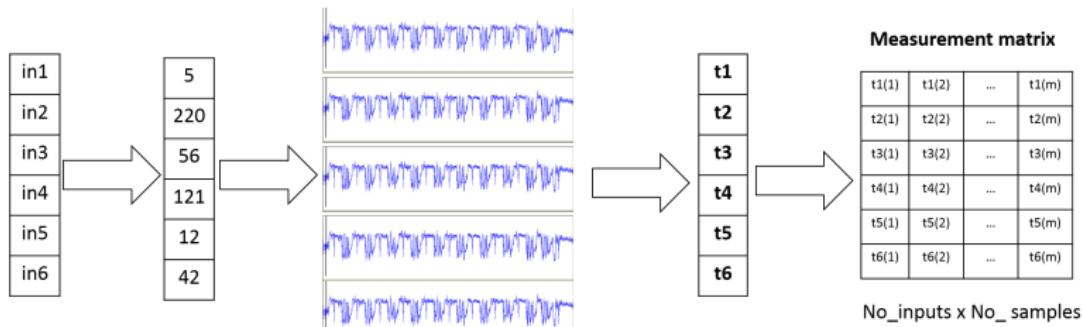
CPA Step 2: Measure the power consumption

- Capturing 6 power traces with m time points (samples) results in the following measurement matrix



CPA Step 2: Measure the power consumption

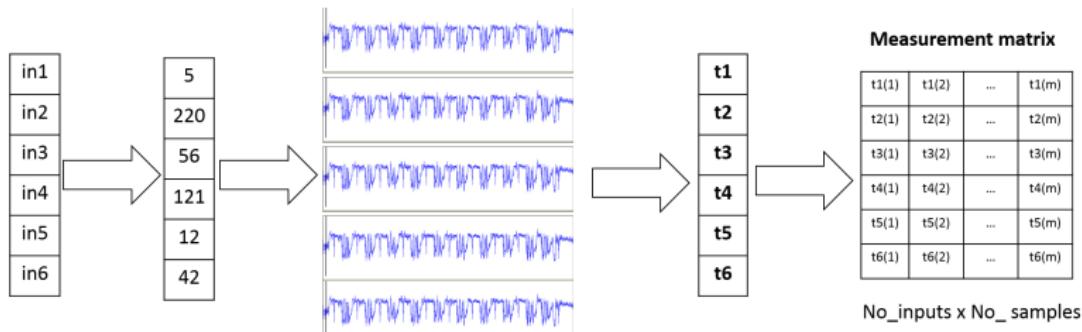
- Capturing 6 power traces with m time points (samples) results in the following measurement matrix



- Note that the power traces originate from the device, i.e. they are related to the secret key stored inside the device

CPA Step 2: Measure the power consumption

- Capturing 6 power traces with m time points (samples) results in the following measurement matrix



- Note that the power traces originate from the device, i.e. they are related to the secret key stored inside the device
- We will refer to the unknown key that is stored in the device as k_{dev}

CPA Step 3: Predict (hypothetical) intermediate values

- In our device $y = Sbox(in \oplus k_{dev})$, but k_{dev} is unknown!
- For a given input in we can compute the value y for all possible keys $k \in \{0, 1, \dots, 255\}$

CPA Step 3: Predict (hypothetical) intermediate values

- In our device $y = Sbox(in \oplus k_{dev})$, but k_{dev} is unknown!
- For a given input in we can compute the value y for all possible keys $k \in \{0, 1, \dots, 255\}$
- ForAll $in \in \mathbf{in}$
 - ForAll $k \in \{0, 1, \dots, 255\}$
 - Compute $y(in, k) = Sbox(in \oplus k)$

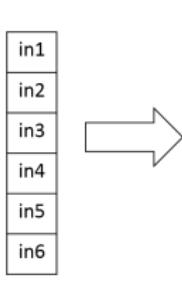
CPA Step 3: Predict (hypothetical) intermediate values

- In our device $y = Sbox(in \oplus k_{dev})$, but k_{dev} is unknown!
- For a given input in we can compute the value y for all possible keys $k \in \{0, 1, \dots, 255\}$
- ForAll $in \in \text{in}$

ForAll $k \in \{0, 1, \dots, 255\}$

Compute $y(in, k) = Sbox(in \oplus k)$

Value-prediction matrix



	k=0	k=1		k=255
Sbox(in1 XOR 0)	Sbox(in1 XOR 1)	...	Sbox(in1 XOR 255)	
Sbox(in2 XOR 0)	Sbox(in2 XOR 1)	...	Sbox(in2 XOR 255)	
Sbox(in3 XOR 0)	Sbox(in3 XOR 1)	...	Sbox(in3 XOR 255)	
Sbox(in4 XOR 0)	Sbox(in4 XOR 1)	...	Sbox(in4 XOR 255)	
Sbox(in5 XOR 0)	Sbox(in5 XOR 1)	...	Sbox(in5 XOR 255)	
Sbox(in6 XOR 0)	Sbox(in6 XOR 1)	...	Sbox(in6 XOR 255)	

No_inputs x No_keys

- One of the columns of this value-prediction matrix is the correct one!
- Divide and Conquer strategy

CPA Step 4: Leakage model

- We map the hypothetical intermediate values to hypothetical power consumption values, producing the power-prediction matrix

Power-prediction matrix

k=0	k=1		k=255
HW(Sbox(in1 XOR 0))	HW(Sbox(in1 XOR 1))	...	HW(Sbox(in1 XOR 255))
HW(Sbox(in2 XOR 0))	HW(Sbox(in2 XOR 1))	...	HW(Sbox(in2 XOR 255))
HW(Sbox(in3 XOR 0))	HW(Sbox(in3 XOR 1))	...	HW(Sbox(in3 XOR 255))
HW(Sbox(in4 XOR 0))	HW(Sbox(in4 XOR 1))	...	HW(Sbox(in4 XOR 255))
HW(Sbox(in5 XOR 0))	HW(Sbox(in5 XOR 1))	...	HW(Sbox(in5 XOR 255))
HW(Sbox(in6 XOR 0))	HW(Sbox(in6 XOR 1))	...	HW(Sbox(in6 XOR 255))

No_inputs x No_keys

- A common choice is Hamming weight but keep in mind that other models may be applicable

CPA Step 5: Comparison

- Compare the hypothetical power consumption values with the real measurements using Pearson correlation

CPA Step 5: Comparison

- Compare the hypothetical power consumption values with the real measurements using Pearson correlation
- ForAll columns of measurement matrix
 - ForAll columns of power prediction matrix
 - Compute the correlation between columns

CPA Step 5: Comparison

- Compare the hypothetical power consumption values with the real measurements using Pearson correlation
- ForAll columns of measurement matrix
 - ForAll columns of power prediction matrix
 - Compute the correlation between columns
- The highest correlation value reveals the key

CPA Step 5: Comparison

- Compare the hypothetical power consumption values with the real measurements using Pearson correlation
- ForAll columns of measurement matrix
 - ForAll columns of power prediction matrix
 - Compute the correlation between columns
- The highest correlation value reveals the key

Pearson correlation coefficient $\rho_k(L, HW(V_g))$ for leakage L , Hamming weight of the intermediate value computed $HW(V)$ and key guess K_g :

$$\rho_k(L, HW(V_g)) = \frac{cov[L, HW(V_g)]}{\sqrt{Var[L] \cdot Var[HW(V_g)]}} \quad (1)$$

CPA Step 5: Comparison

